

Spotlight Reflex

Minimal-Shot Autonomous Driving via Grounded ManeuverTokens

Alba Maria Tellez Fernandez

May 2026

amtellezfernandez@gmail.com · Waymo Open Dataset E2E Challenge

Who I Am & What I Built

Alba Maria Tellez Fernandez

AI / robotics builder focused on autonomy systems that can operate outside heavily curated training distributions.

Question behind this project:

- Can a compact geometric representation replace memorized scene priors?
- Does “*avoid occupied space, preserve route*” survive new geometry?
- What breaks when the same policy enters AlpaSim?

Built from scratch:

- custom long-tail simulator + evaluation suites
- Spotlight Reflex token policy runtime
- DAgger, scorer, and hybrid policy experiments
- AlpaSim bridge for WOD-E2E scenes
- WOD-E2E selector with Cosmos features

Main Claim

Minimal-shot autonomy needs **grounded action interfaces**: bounded maneuver candidates, explicit route/occupancy geometry, and transfer metrics separated by axis.

What this is

- a runnable simulator and transfer harness
- grounded token-selection experiments
- reproducible artifacts and audit paths

What this is not

- not a production AV stack
- not a LaST-VLA reproduction
- not a final Waymo hidden-test claim

The Problem — Continuous Policy Collapse in Long-Tail

Imitation learning minimises:

$$\min_f \mathbb{E}_{(x,y) \sim \mathcal{D}} [\mathcal{L}(f(x), y)]$$

Works on the common driving distribution $\mathcal{D}$.

Fails when $P_{\text{test}} \neq P_{\text{train}}$: continuous distributions spread — the policy assigns probability mass to unsafe intermediate states that never appeared in training.

WOD-E2E targets **11 scenario types**

occurring $<0.03\%$ of driving time: wrong-way actors, animals, debris, construction, cut-ins, cyclists.

Continuous baseline: no world-state reasoning.

Drives directly into the actor. Gauntlet pass rate:

2.1%, collision: **20.5%**.



Wrong-way actor, night. Continuous baseline: no awareness the path is occupied. Policy variance under OOD $\rightarrow$ collision.

Motivation: Grounded Decisions, Not Language Plans


Recent VLA work gives a useful direction: driving decisions need physical grounding, because language is a lossy interface for meter-level geometry, timing, clearance, and acceleration.

Related context: LaST-VLA (arXiv:2603.01928)

Xiaomi EV + Tsinghua align latent states with **world-model dynamics** and **3D geometric priors**.

That is a foundation-model path to physical grounding.

That leaves our complementary question: after actions are grounded, **what transfers across simulators?** Our stack exposes route blockage, clearance, actor motion, and fixed maneuver tokens, so failures separate into collision, lane, route, and progress axes.



Positioning: LaST-VLA is a large latent VLA planner. This project is not a smaller copy: it is a minimal-shot transfer stack that exposes when grounded token selection helps, when it breaks, and which metric axis caused the break.

Three Claims

Claim 1 — Discrete Bottleneck Hypothesis

Quantizing the action space to 9

ManeuverTokens is a structural robustness prior.

Fixed geometric semantics prevent policy variance from spreading to unsafe intermediate states.

Evidence: 0 collisions / 350 rollouts; 27× gauntlet pass.

Claim 2 — Geometric Grounding of Visual Latents

Cosmos 64d embeddings carry preference signal inaccessible to linear models but extractable by a GPU MLP. Generation-oriented VLAs fail on discrimination tasks.

Evidence: linear head $\rightarrow$ no gain; GPU MLP $\rightarrow$ +0.042 RFS.

Claim 3 — Causal Bottleneck Analysis

ManeuverToken selection is causally determined

by geometric pressure scalars — no demonstration data, no spurious correlations.

This precisely localizes the 1.419 RFS oracle gap to the visual-to-geometric bridge.

Evidence: failure modes trace to scalar interactions; oracle gap not in action space or ranker.

What I Built: Simulator First, AlpaSim Second

1. Controlled simulator

I built a lightweight closed-loop 2D AV simulator because it gives exact ground truth and controllable long-tail stress tests.

- 11 WOD-inspired scenario clusters
- compositional OOD, hidden, adversarial, gauntlet suites
- paired seeds and continuous metrics: collision, clearance, progress, jerk

Purpose: isolate whether the policy really follows route/occupancy geometry when the world state is known.

2. AlpaSim transfer bridge

Then I ported the **same token-policy API** into AlpaSim, using WOD-E2E front-camera scenes and a route/ego adapter.

- camera + route + ego dynamics → proxy geometry
- same tokens, same policy call, different execution stack
- per-frame logs for token choice and transfer failure axes

Purpose: test whether the internal simulator result survives sensor-realistic execution and proxy-state reconstruction.



Core Idea: Reason from Geometry, Not Memory

*"A trajectory that avoids occupied space and respects route geometry is correct **regardless of what the obstacle is called.**"*

Formal invariance: let T be any relabelling of object identities, scenario names, or cluster tags that leaves occupancy G and route R unchanged.

$$M(s) = M(T(s)) \quad \forall s, \forall T$$

"sheep" $\rightarrow$ "debris" changes nothing.

Regression-tested: relabel all objects in every scene and assert maneuver, score, and reason stay identical.



*Same wrong-way actor, night. Spotlight Reflex selects *evasive_right*, clears the actor, and returns to lane. Geometry resolved the decision.*

Architecture: Six World-State Scalars

All signals computed from raw obstacle positions and route geometry. **Nothing is learned; nothing is scene-specific.**

Signal	Definition	Range
obstacle_pressure	$p = \text{clip}\left(\frac{10-d_{\min}}{10}, 0, 1\right)$	$[0, 1]$
route_blockage	fraction of forward corridor occupied	$[0, 1]$
corridor_blocked	hard blockage within stopping distance	bool
left_clearance	free lateral space to the left (m)	≥ 0
right_clearance	free lateral space to the right (m)	≥ 0
preferred_escape_side	$\arg \max(\text{left}, \text{right})$	L / R

$d_{\min}$ = distance to nearest obstacle. Pressure = 1.0 at contact, 0 beyond 10 m. No object identity, no semantic labels, no history. All O(1) to compute. Formal invariance tested: $M(s) = M(T(s))$ for any object-identity relabelling T .

Architecture: Nine ManeuverTokens — The Discrete Bottleneck

Each token: **20-point, 5 s trajectory at 4 Hz** via Hermite smoothstep. Evasive maneuvers apply the lateral offset in the first third for fast separation.

Maneuver	Intent	Speed	Lateral
stop	Emergency stop	0.00×	0.0 m
crawl	Very slow progress	0.15×	0.0 m
maintain	Hold current speed	1.00×	0.0 m
slow_yield	Decelerate & yield	0.50×	0.0 m
nudge_left	Minor left shift	0.80×	+0.9 m
nudge_right	Minor right shift	0.80×	−0.9 m
evasive_left	Strong left evasion	0.70×	+2.2 m
evasive_right	Strong right evasion	0.70×	−2.2 m
lane_recover	Return to centre	0.90×	0.0 m

Why discrete?

A continuous policy assigns mass to unsafe *intermediate states* between tokens. The token constraint eliminates those states by construction — uncertainty changes **scores**, not physical footprints.

Coverage: 5 longitudinal × 4 lateral directions covers all safety-critical responses in the Waymo long-tail taxonomy. Nothing missing; nothing extra.

Architecture: Trust-Region Scoring & Decision Flow

Speed-scaled tolerance:

$$s(v) = \text{clip}\left(\frac{v-1.4}{9.6}, 0, 1\right) \times 0.5 + 0.5$$

Horizon	Lateral	Longitudinal
3 s	$1.0 \times s(v) \text{ m}$	$4.0 \times s(v) \text{ m}$
5 s	$1.8 \times s(v) \text{ m}$	$7.2 \times s(v) \text{ m}$

$$S(\tau) = 0.5 \cdot \max_r \text{score}_{3s} + 0.5 \cdot \max_r \text{score}_{5s}$$

Scores ≤ 50 = unsafe floor. Clearance
 $< 0.55 \text{ m}$: -1000 penalty.

Decision pipeline (4 steps):

1. Scene $\rightarrow$ compute 6 world-state scalars
2. Reference rule engine maps scalars to expected tokens + scores (76–94)
3. Each of the 9 tokens scored against references at 3 s and 5 s
4. Token with highest combined score is selected

Every step emits a full log: maneuver, scores, decision reason, all scalar values. In AlpaSim: serialised as `reasoning_text` JSON. **Identical policy binary** in 2D simulator and AlpaSim.

Policy Architecture — Decision Flow



6 Scalars: pressure, blockage, corridor blocked, left/right clearance, preferred escape side.

Ref Rules: world-state thresholds → expected maneuver + score. Scores ≤ 50 are unsafe.

Scorer: 3 s + 5 s horizons, speed-scaled tolerance, full log emitted every step.

Simulation: Gauntlet — The Bottleneck Test

Gauntlet: 4 simultaneous hazards, 3.5 m corridor, 420 matched scenarios (seeds 1–80).

Policy	Pass	Collision
Continuous baseline	2.1%	20.5%
Spotlight Reflex	57.6%	7.9%
<i>Improvement</i>	27×	2.6× fewer

Token constraint eliminates unsafe intermediate states. Baseline spreads mass across trajectory space; collision rate: **20.5%**.



*evasive_right, clears wrong-way actor,
recovers lane centre.*

Simulation: 350 Rollouts — Zero Collisions

Primary benchmark — seeds 1–10, deterministic. Same policy binary across all suites.


Suite	Runs	Coll.	Pass
WOD-style (11 clusters)	110	0	100%
Compositional OOD	60	0	100%
Adversarial (2–3 hazards)	60	0	100%
Hidden holdout	60	0	100%
Gauntlet (4h, 3.5 m)	60	0	60%
Total	350	0	93.1%

COMPASS **9.137/10** over 700 ranked runs

95% CI collision rate: **[0.0, 0.0053]**


Mean minimum clearance: 2.96 m

OOD robustness: compositional generator samples
topology $\times$ hazard $\times$ weather independently. 8
topologies $\times$ 11 hazards $\times$ 4 weather modes.



Construction zone — 5.2 m corridor, cone field, lane closure. Token: `nudge_right`. Policy selects from geometry, not scenario label.

AlpaSim: What We Built




Real Waymo front-camera frames alongside adapter output. The policy API is fixed; the adapter reconstructs proxy geometry.

Engineering work completed:

- Local AlpaSim plugin wrapping the same token policy API
- Front-camera + route + ego-dynamics adapter
- Proxy-state reconstruction into Spotlight/DAGger features
- Docker/runtime bootstrap scripts
- Scene presets and paired matrix runner
- Per-frame token, adapter, collision, offroad, lane, progress logs

This made the simulator result auditable in a second execution stack, instead of staying as a closed internal demo.

AlpaSim: External Transfer Matrix



Same AlpaSim bridge, now used to run paired learned-policy transfer matrices over shared WOD-E2E clips.

10 shared WOD-E2E clips, paired by scene:

Variant	Coll.	Offroad	Lane	Prog.
Raw DAgger iter2	0.60	0.90	0.70	0.034
Clamped DAgger	0.70	0.50	0.20	0.384
Hard-veto hybrid	0.80	0.20	0.70	0.837
Source decay	0.60	0.90	0.40	0.175

Interpretation: each intervention fixes a different axis. Clamping improves route geometry; hard-veto recovers progress/offroad; source decay helps lane behavior. No single learned variant dominates all axes.

This is the transfer finding: internal success is not enough unless the selector is grounded at execution time.

WOD-E2E Benchmark: Trajectory Selector

479 preference-labeled Waymo frames · segment-grouped 5-fold CV · RFS = preference alignment

Selector	Folds	RFS
Official Waymo baseline	—	7.022
Local CV baseline	—	7.131
Gate only (HGB)	5	7.803
RFF direct policy	5	7.834
GPU MLP + Cosmos 64d	5	7.845
HGB Optuna peak [†]	2	7.880
Oracle (perfect)	5	9.264

[†] 2-fold; not comparable to 5-fold results.

Pipeline: kinematic + ridge + temporal candidates
→ HGB gate (137 features) → GPU MLP override
(4.2% of frames, precision 0.60).

Visual embedding results:

Config	Head	RFS
No embedding	Ridge	7.803
Cosmos 64d	Ridge	≈7.803
InternVLA 128d	Ridge	≈7.803
RFF gate	non-lin.	7.834
Cosmos 64d	GPU MLP	7.845

Cosmos embeddings carry signal — only a non-linear model extracts it. InternVLA fails: generation $\neq$ discrimination.

GPU MLP — h=64, lr=0.003866, batch=512, 10 epochs, Optuna TPE.

Oracle Gap: Causal Localization of the Bottleneck

Oracle gap: 1.419 RFS (9.264 – 7.845)

The right candidate exists in the pool. Where does the gap live?

Slice	Sel.	Regret
GO_STRAIGHT ($n=427$)	7.959	1.377
GO_LEFT ($n=23$)	7.107	1.614
GO_RIGHT ($n=29$)	6.574	2.063
speed:urban ($n=136$)	8.132	1.285
speed:slow ($n=133$)	7.564	1.632
speed:fast ($n=44$)	7.957	1.479

Causal elimination:

- **Not the action space** — oracle 9.264; preferred candidate exists in pool.
- **Not the ranker** — gate 7.803 + GPU MLP 7.845 exhaust trajectory + visual feature information.
- **The missing bridge:** model that infers geometric context from camera would resolve the $\sim 15\%$ of frames where gate and direct policy both fail.

GO_RIGHT (2.063 regret, $n=29$): temporal gate over-selects. Rightward intent requires lateral clearance checks absent from ego-history — a geometry-aware camera model would catch this.

Three Contributions — Backed by Evidence

Contribution	Evidence
1. Discrete Bottleneck	27× gauntlet pass (57.6% vs 2.1%). 0 collisions / 350 rollouts. COMPASS 9.137/10. Collision: 7.9% vs 20.5% baseline under 4 hazards.
2. Geometric Grounding	Linear head on Cosmos 64d: no gain. GPU MLP: +0.042 RFS. InternVLA: no gain (generation $\neq$ discrimination).
3. Causal Localization	Failure modes trace to scalar interactions, not object labels. 1.419 RFS oracle gap $\rightarrow$ visual-geometric bridge missing.
AlpaSim bridge	Same token-policy API runs through a front-camera/route/ego adapter with paired scene matrices and per-frame axis logs.
AlpaSim finding	10 paired WOD-E2E clips expose axis-specific transfer: clamping helps route geometry, veto helps progress, source decay helps lane behavior.
WOD-E2E champion	7.845 RFS (5-fold), +0.823 vs official Waymo baseline.

Conclusion & Next Steps

What we showed:

- **Discrete bottleneck:** 9 tokens, 0 collisions, $27\times$ gauntlet gain.
- **Grounded visual signal:** Cosmos helps only with a nonlinear readout.
- **Bottleneck localized:** 1.419 RFS oracle gap points to the visual-to-geometric bridge.
- **AlpaSim bridge:** same policy API, second execution stack, paired WOD-E2E scenes, per-axis logs.

General principle:

Grounded action interfaces make transfer failures measurable instead of hidden inside one aggregate score.

Next steps:

1. **Preference-aligned visual encoder**
Fine-tune on triplets (frame, win, loss) to close the oracle gap.
2. **Intent calibration**
Stratified reweighting by intent; GO_RIGHT regret is 2.063.
3. **Test-set run**
Pipeline complete; missing gated Waymo test TFRecords.
4. **Physical deployment path**
Obstacle sensors + route command → Spotlight Reflex.

Alba Maria Tellez Fernandez

amtellezfernandez@gmail.com

`docs/simulation.md · docs/wod-e2e-system-walkthrough.md`

`docs/corl2027/paper.pdf · docs/corl2027/paper.txt`